

I run all the programs on EECS hydra clusters.

## MPI Datatype

diagonal : diagonal.c

a datatype describing matrix diagonal.

Command: `mpirun -np 2 diagonal` (The process number is fixed)

Input:

```
1, 2, 3, 4, 5, 6, 7, 8,
9, 10, 11, 12, 13, 14, 15, 16,
17, 18, 19, 20, 21, 22, 23, 24,
25, 26, 27, 28, 29, 30, 31, 32,
33, 34, 35, 36, 37, 38, 39, 40,
41, 42, 43, 44, 45, 46, 47, 48,
49, 50, 51, 52, 53, 54, 55, 56,
57, 58, 59, 60, 61, 62, 63, 64
```

Output: rank= 1 b= 1 10 19 28 37 46 55 64

C arrays are stored in row-major order while FORTRAN arrays are stored in column-major order.

transpose : transpose.c

a datatype for matrix transpose

Command: `mpirun -np 2 transpose` (The process number is fixed)

Input:

```
1, 2, 3, 4, 5, 6, 7, 8,
9, 10, 11, 12, 13, 14, 15, 16,
17, 18, 19, 20, 21, 22, 23, 24,
25, 26, 27, 28, 29, 30, 31, 32,
33, 34, 35, 36, 37, 38, 39, 40,
41, 42, 43, 44, 45, 46, 47, 48,
49, 50, 51, 52, 53, 54, 55, 56,
57, 58, 59, 60, 61, 62, 63, 64
```

Output:

```
rank= 1
1 9 17 25 33 41 49 57
2 10 18 26 34 42 50 58
3 11 19 27 35 43 51 59
4 12 20 28 36 44 52 60
```

```
5 13 21 29 37 45 53 61
6 14 22 30 38 46 54 62
7 15 23 31 39 47 55 63
8 16 24 32 40 48 56 64
```

## Ring Application

ring : ring.c

a simple application that passes a token between all the processes in one communicator.

Command: `mpirun -np 4 ring`

Output:

```
rank= 1, received token 0.
rank= 2, received token 1.
rank= 3, received token 2.
```

## Process Grids & Multi-pipeline communications

I put these two tasks together. I first create the row and column communicators and then do the iterations to pass tokens.

grids : grids.c

Command: `mpirun -np 9 grids` (The process number is fixed)

Input: These are the original tokens.

```
1, 2, 3, 4, 5, 6,
7, 8, 9, 10, 11, 12,
13, 14, 15, 16, 17, 18,

19, 20, 21, 22, 23, 24,
25, 26, 27, 28, 29, 30,
31, 32, 33, 34, 35, 36
```

$P=3$ ,  $Q=3$ , so in this example above, process 0 has 1, 4, 19, 22. The tokens will be passed through rows and columns and back to their original position in the end. I only output the first token and second token on the process 0 after the iterations. If they are the same with the original one, then the program is correct.

$P*Q$  is the grid size. So we need  $P+Q-2$  steps to have a full exchange of tokens.

Output:

```
Row iteration finishes.
rank[0], 1st token: 1, 2nd token: 4
Column iteration finishes.
rank[0], 1st token: 1, 2nd token: 4
```

## PDGEMM

I only implemented a simple matrix matrix multiplication.

`PDGEMM("No transpose", "No transpose", N, N, N,  $\alpha = 1$ , A, B,  $\beta = 0$ , C)`

And I haven't successfully implemented the task with all the requirements, e.g. a user supplied argument k. But this parallel version works and the result is accurate.

Command: `mpirun -np 6 a1.txt b1.txt` (a1.txt and b1.txt are filenames)

Output: (a matrix C)

Elapsed time: 0.000057 seconds.

The Matrix C.

2.119985 2.773072 1.938637 2.246718 1.949592 3.105905 2.937909 2.226654 2.549402 ...

..... (The detailed output of the matrix is omitted here.)